

Reinforcement learning and robot navigation

Charles Dufour


May 2, 2018



The problem


- Framework: a raspberry pie 3 robot which can follow lines
- The task: the robot should adapt its speed with respect to traffic lights
- How: using Reinforcement Learning (RL) and Markov Decision Process (MDP)

The task




- Part I: theoretical background
- Part II: results from first implementations

Reinforcement learning



The agent's job is to find a behavior that maximizes the long-run sum of values of the rewards.

Markov Decision Process intuition



actions possible in state 0:

① \rightarrow : action 1

② \rightarrow : action 2

Episode : $s_0; a_1, s_1, r_1; a_2, s_2, r_2; \dots; a_n, s_n, r_n$

Definition: (MDP)

- A set of states $\mathcal{S} = \{s_0, s_1, s_2, \dots, s_n\}$
- A set of actions $\mathcal{A} = \{a_1, a_2, a_3, \dots, a_k\}$
- A transition function $T(a, s, s') = \mathbb{P}[s' \mid a, s]$
- A reward function $R : \mathcal{S} \mapsto \mathbb{R}$
- A discount factor $0 \leq \gamma < 1$

Markov property

The transitions only depends on the current state and the current action.

Definition: (policy)

A *policy* π is a probabilistic mapping from the set of states to the set of actions :

$$\pi : \mathcal{A} \times \mathcal{S} \mapsto [0, 1]$$

$$\text{s.t. } \sum_a \pi(a \mid s) = 1 \quad \forall s \in \mathcal{S}$$

How to

How to assess the quality of policies so we can find the best one?
What is the best policy?

How to assess the quality of policies

We are interested in maximizing the discounted return: G_t

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} = R_{t+1} + \gamma G_{t+1}$$

How to assess the quality of policies

We are interested in maximizing the discounted return: G_t

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} = R_{t+1} + \gamma G_{t+1}$$

Action value while in a state s under π

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a]$$

How to assess the quality of policies

We are interested in maximizing the discounted return: G_t

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} = R_{t+1} + \gamma G_{t+1}$$

Action value while in a state s under π

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a]$$

State value under policy π

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t \mid S_t = s]$$

Recursive definition of value functions

action value

$$q_{\pi}(s, a) = \sum_{s'} \mathbb{P}(s' \mid a, s) [R(s') + \gamma v_{\pi}(s')]$$

state value

$$v_{\pi}(s) = \sum_a \pi(a \mid s) \sum_{s'} \mathbb{P}(s' \mid a, s) [R(s') + \gamma v_{\pi}(s')]$$

How to assess the quality of policies

How to compare two policies

$$\pi \leq \pi' \iff v_{\pi}(s) \leq v_{\pi'}(s) \quad \forall s \in \mathcal{S}$$

How to assess the quality of policies

How to compare two policies

$$\pi \leq \pi' \iff v_{\pi}(s) \leq v_{\pi'}(s) \quad \forall s \in \mathcal{S}$$

Optimal policy

$$\pi_* \quad s.t. \quad \forall \pi : \pi_* \geq \pi$$

Bellman optimality equations

We can associate the value function v_* and q_* to the optimal policy π_*

$$v_*(s) = \max_a q_*(s, a)$$

$$q_*(s, a) = \sum_{s'} \mathbb{P}(s' \mid s, a) [R(s') + \gamma \max_{a'} q_*(s', a')]$$

Finding optimal policy from value functions

$$\pi_*(s) = \operatorname{argmax}_a q_*(s, a)$$

Computational issue

- | \mathcal{S} | linear equations to solve to evaluate policy
- | \mathcal{S} | non linear equations to solve the Bellman optimality equation

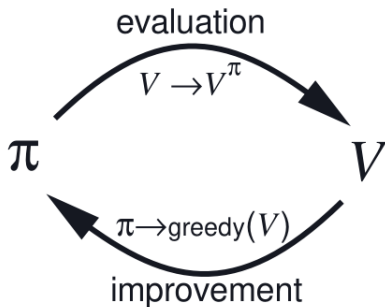
Computational issue

- | \mathcal{S} | linear equations to solve to evaluate policy
- | \mathcal{S} | non linear equations to solve the Bellman optimality equation

Approximation of value function
and
Policy iteration

Policy iteration algorithm


solving MDP using dynamic programming




1

¹From (Sutton & Barto, 1998)

Our problem



First a simpler problem



states

- position $\{0,1,2,\dots,L, \text{Lava}\}$
- speed $\{\text{low}, \text{medium}, \text{high}\}$

states

- position $\{0,1,2,\dots,L, \text{Lava}\}$
- speed $\{\text{low}, \text{medium}, \text{high}\}$

actions

- slowing down
- maintaining speed
- speeding up

states

- position $\{0,1,2,\dots,L, \text{Lava}\}$
- speed $\{\text{low, medium, high}\}$

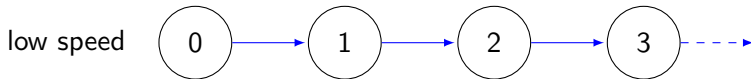
actions

- slowing down
- maintaining speed
- speeding up


reward function

- Lava: reward of $-L$
- L in low speed: reward of $+L$
- any other state : reward of -1

Maintaining speed

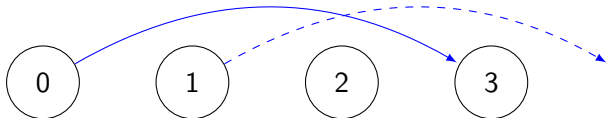


Maintaining speed

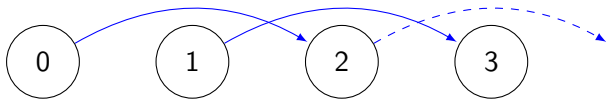


Maintaining speed

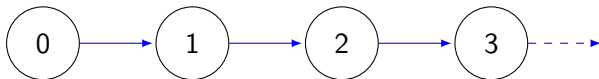
high speed



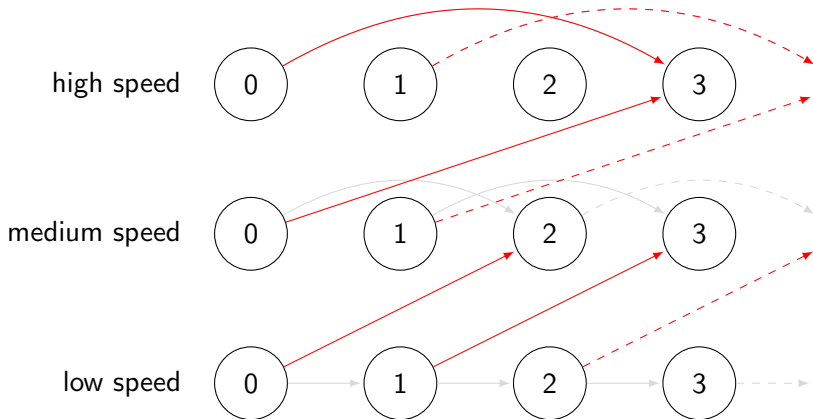
medium speed



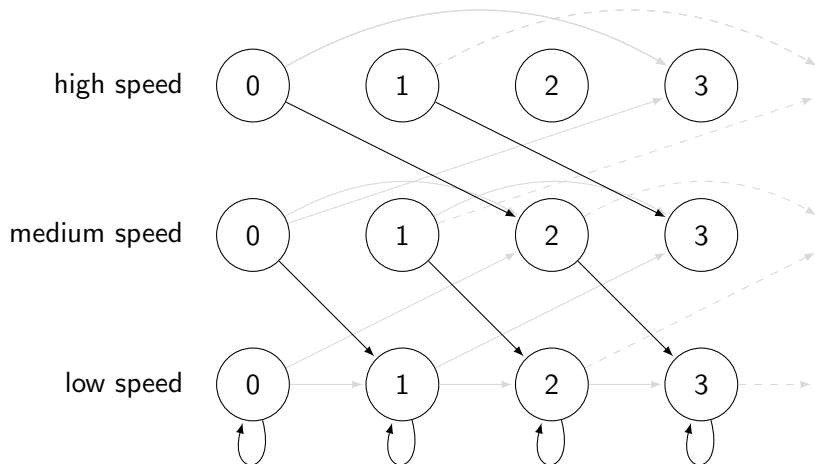
low speed



Speeding up




Slowing down



State value function at different iterations

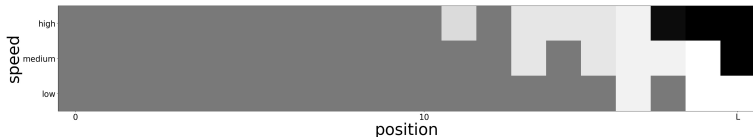
the brighter the better the value



States values at iteration 0, all the state values are equal to 0

State value function at different iterations

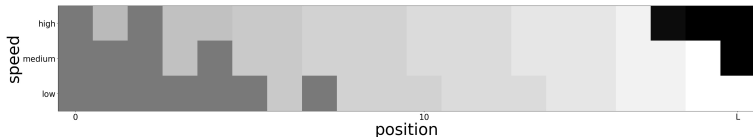
the brighter the better the value



States values at iteration 2

State value function at different iterations

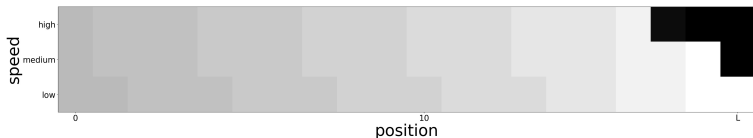
the brighter the better the value



States values at iteration 4

State value function at different iterations


the brighter the better the value



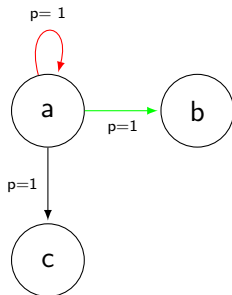
States values at iteration 6 (where stable policy is attained)

Results for deterministic actions

One of the optimal trajectories

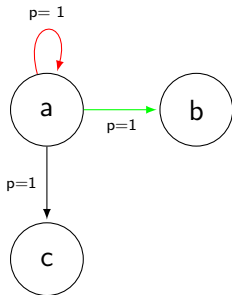


Deterministic actions are not really realistic

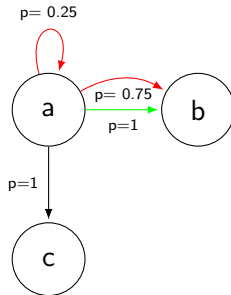


Deterministic actions

Deterministic actions are not really realistic

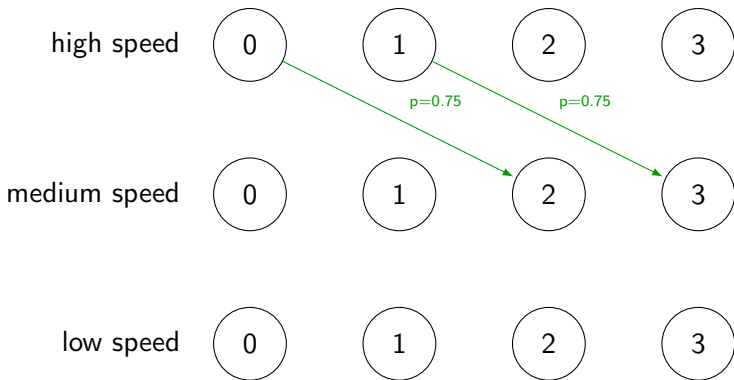


Deterministic actions

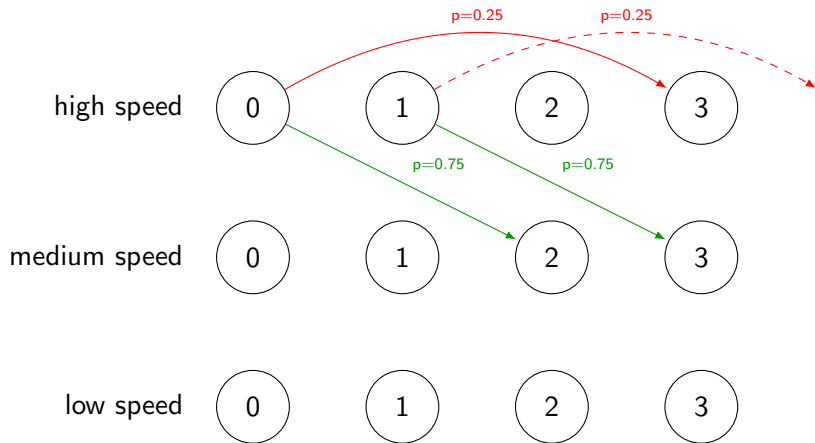


Stochastic actions

Decelerating for decelerating being a stochastic action

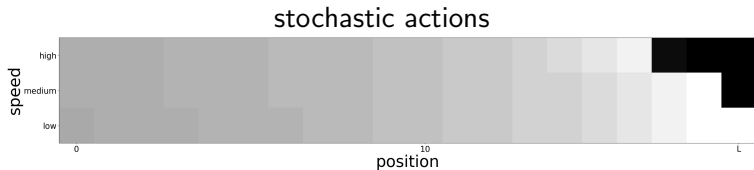


Decelerating for decelerating being a stochastic action




Results for stochastic actions

state-value function at the end of the iterations



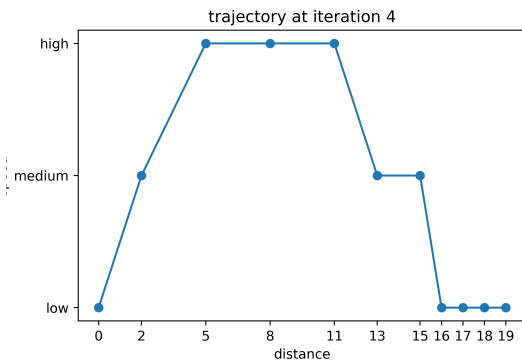
Results for stochastic actions

decelerating results in keeping the same speed with probability $1/4$



Results for stochastic actions

decelerating results results in keeping the same speed with probability $1/4$



where stable policy is attained

- We introduced the theoretical framework: Reinforcement learning and Markov Decision Processes
- We solved a simple problem where we know the model (the transition function in particular)

What's next ?

already working on

- Add the traffic light into this setting
- Find the distance from the robot's camera to the traffic light
- getting data from the robot

In a not so distant future

- Explore other algorithms and compare them : Monte-Carlo methods, temporal difference learning, Q learning, Expected Sarsa...
- Neuro-dynamic programming ?

What's next ?

already working on

- Add the traffic light into this setting
- Find the distance from the robot's camera to the traffic light
- getting data from the robot

In a not so distant future

- Explore other algorithms and compare them : Monte-Carlo methods, temporal difference learning, Q learning, Expected Sarsa...
- Neuro-dynamic programming ?

Ultimately

Implement everything on the robot ...

What's next ?

already working on

- Add the traffic light into this setting
- Find the distance from the robot's camera to the traffic light
- getting data from the robot

In a not so distant future

- Explore other algorithms and compare them : Monte-Carlo methods, temporal difference learning, Q learning, Expected Sarsa...
- Neuro-dynamic programming ?

Ultimately

Implement everything on the robot ... and pray that everything works well